

COMP 9517 Computer Vision

Segmentation and Feature Tracking




Curves and Segmentation

- curves corresponding to object boundaries are common, especially in natural environment
- one approach to segmentation is to locate object boundary curves in images





Active Contours

- Aim: to locate boundary curves in images
- How: boundary detectors iteratively move towards their final solution under a combination of *image*, *smoothness* and *optional user-guidance* forces



Active Contours

- Examples of implementations:
 - Snakes
 - Level sets




- Active contours can also be used in a wide variety of object-tracking applications

Global Optimization Techniques

- formulate the goals of the desired transformation (eg segmentation) using some optimisation criterion, then infer the solution that best meets this criterion
- Problem: finding a smooth surface that passes through a set of measured data points
 - ill-posed: many possible surfaces can fit this data
 - ill-conditioned: small changes in input can sometimes lead to large changes in the fit
 - inverse problems: recover the unknown function from which data points were sampled
- therefore, regularisation is needed
 - to fit models to data that badly underconstrain the solution space

Global Optimisation

- in order to quantify what it means to find a smooth solution, we can define a norm on the solution space
- the **derivative** is a measure of how a function changes as its input changes.



- a derivative can be thought of as how much one quantity is changing in response to changes in some other quantity
- for example, the derivative of the position of a moving object with respect to time is the object's instantaneous velocity

$$m = \frac{\Delta f(x)}{\Delta x} = \frac{f(x+h) - f(x)}{(x+h) - (x)} = \frac{f(x+h) - f(x)}{h}.$$

Global Optimisation

- For one dimensional function $f(x)$, to find a smooth solution, we can define a **norm** on the solution space:
 - integrate the squared first derivative of the function

$$\varepsilon_1 = \int f_x^2(x) dx$$

- integrate the squared second derivative

$$\varepsilon_2 = \int f_{xx}^2(x) dx$$

- For two dimensions (e.g. images), the corresponding smoothness functions are partial derivatives:

$$\varepsilon_1 = \int f_x^2(x, y) + f_y^2(x, y) dx dy = \int \|\nabla f(x, y)\|^2 dx dy$$

$$\varepsilon_2 = \int f_{xx}^2(x, y) + 2f_{xy}^2(x, y) + f_{yy}^2(x, y) dx dy$$

Global Optimization

- We also require a data term (or data penalty)
- For scattered data interpolation, the data term measures distance between function $f(x,y)$ and a set of data points $d_i = d(x_i, y_i)$
- By adding the norm and the data term, we get a global energy term that can be minimized

Snakes

- Snakes are a 2-D generalisation of the 1-D energy minimising splines:

$$\mathcal{E}_{\text{int}} = \int \alpha(s) \|f_s(s)\|^2 + \beta(s) \|f_{ss}(s)\|^2 ds$$

- discretised version of the energy function:

$$E_{\text{int}} = \int \alpha(s) \|f(i+1) - f(i)\|^2 / h^2 + \beta(s) \|f(i+1) - 2f(i) + f(i-1)\|^2 / h^4$$

- additional external image-based and constraint-based potentials:

$$\mathcal{E}_{\text{image}} = \omega_{\text{line}} \mathcal{E}_{\text{line}} + \omega_{\text{edge}} \mathcal{E}_{\text{edge}} + \omega_{\text{term}} \mathcal{E}_{\text{term}}$$

Snakes

- In practice, only the edge term is used
 - directly proportional to the image gradients:

$$E_{edge} = \sum_i -\|\nabla I(f(i))\|^2$$

- or a smoothed version of the image Laplacian:

$$E_{edge} = \sum_i -\|(G_\delta \circ \nabla^2 I)(f(i))\|^2$$

- distance map to the edges
- user-placed constraints, for e.g.


$$E_{spring} = k_i \|f(i) - d(i)\|^2$$

Snakes

- Put them together:

$$E = E_{\text{int}} + E_{\text{image}} + E_{\text{spring}}$$


- To minimize E , let E be zero, then energy E can be iteratively minimised by moving the curve



(a)



(b)



Dynamic Snakes


- When the object of interest is being tracked from frame to frame, prediction of new estimates is needed
- How? Using dynamic models:
 - Kalman Filtering
 - Particle Filtering

These will be covered when studying Motion

Level Set Methods


- Limitations of active contours based on parametric curves:
 - challenging to change the topology of the curve as it evolves
 - curve reparameterisation may be required when the shape changes dramatically
- Use the zero crossing of a signed distance function is used to define the curve:

$$\frac{\partial \varphi}{\partial t} = v |\nabla \varphi|.$$




Level Set Methods

- Recent approaches recast the problem in a segmentation framework, where the energy measures the consistency of the image statistics inside and outside the segmented regions




(a)



Mean Shift and Model Finding

- How would you segment this image based on colour alone?



Mean Shift and Model Finding

- **K-means and mixtures of Gaussians** technique uses a parametric model of the density function to answer this question
 - assume the density is the superposition of a small number of simple distributions (e.g., Gaussians) whose locations (centres) and shape (covariance) can be estimated
- **Mean shift** smoothes the distribution and finds its peaks, as well as the regions of feature space that correspond to each peak (non-parametric)

K-mean Clustering


- K-means
 - models the probability density as a superposition of spherically symmetric distributions
 - assumes we are given the number of clusters k
 - randomly initialises sampling centres
 - iteratively updates the cluster centre location based on the samples that are closest to each centre
 - techniques exist for splitting or merging cluster centres to accelerate the process

K-means Clustering

- *Iterative K-means clustering*
 - set iteration count ic to 1
 - randomly choose a set of K means $m_1(1), m_2(1), \dots, m_K(1)$
 - for each vector x_i compute distance $D(x_i, m_k(ic))$ for each $k = 1, \dots, K$ and assign x_i to the cluster C_j with the nearest mean
 - increment ic by 1 and update the means to get a new set $m_1(ic), m_2(ic), \dots, m_K(ic)$
 - Repeat step 3 and 4 until $C_k(ic) = C_k(ic + 1)$ for all k

K-means Clustering

- *Iterative K-means clustering*



K-means and Mixtures of Gaussians

- beside means, each cluster centre is augmented by a covariance matrix re-estimated from the corresponding samples
- instead of using nearest neighbours to associate input samples with cluster centres, a Mahalanobis distance is used:
- samples can be associated with the nearest cluster centre, in *hard assignment* of membership or, *softly assigned* to several nearby clusters
- corresponds to iteratively re-estimating the parameters for a mixture of Gaussians density function:

$$p(x | \{x_i, m_k, S_k\}) = \prod_k p_k N(x | m_k, S_k)$$

$$N(x | m_k, S_k) = \frac{1}{|S_k|} e^{-d(x_i, m_k; S_k)}$$

K-means and Mixture of Gaussians

Expectation Maximisation (EM) can be used to iteratively compute maximum likely estimate for the unknown mixture parameters

The expectation stage (E step) estimates the responsibilities:

$$z_{ik} = \frac{1}{Z_i} p_k N(x_i | m_k, S_k) \quad \sum_k z_{ik} = 1$$

How likely a sample was generated from cluster k

The maximisation stage (M step) updates the parameter values

$$\begin{aligned} \mu_k &= \frac{1}{N_k} \sum_i z_{ik} x_i \\ \Sigma_k &= \frac{1}{N_k} \sum_i z_{ik} (x_i - \mu_k)(x_i - \mu_k)^T \\ \pi_k &= \frac{N_k}{N} \\ N_k &= \sum_i z_{ik} \end{aligned}$$

No of samples in cluster k

Mean Shift

- procedure for locating the maxima of a density function given *discrete data* sampled from that function
- uses a smooth continuous non-parametric model to model the probability density function being segmented
- efficiently finds peaks in this high-dimensional data distribution, without ever computing the complete function explicitly

Mean Shift

- Smooth the data by convolving with a fixed kernel function (kernel density estimation)

$$f(x) = \sum_i K(x - x_i) = \sum_i k\left(\frac{\|x - x_i\|^2}{h^2}\right)$$

- Calculate the gradient:

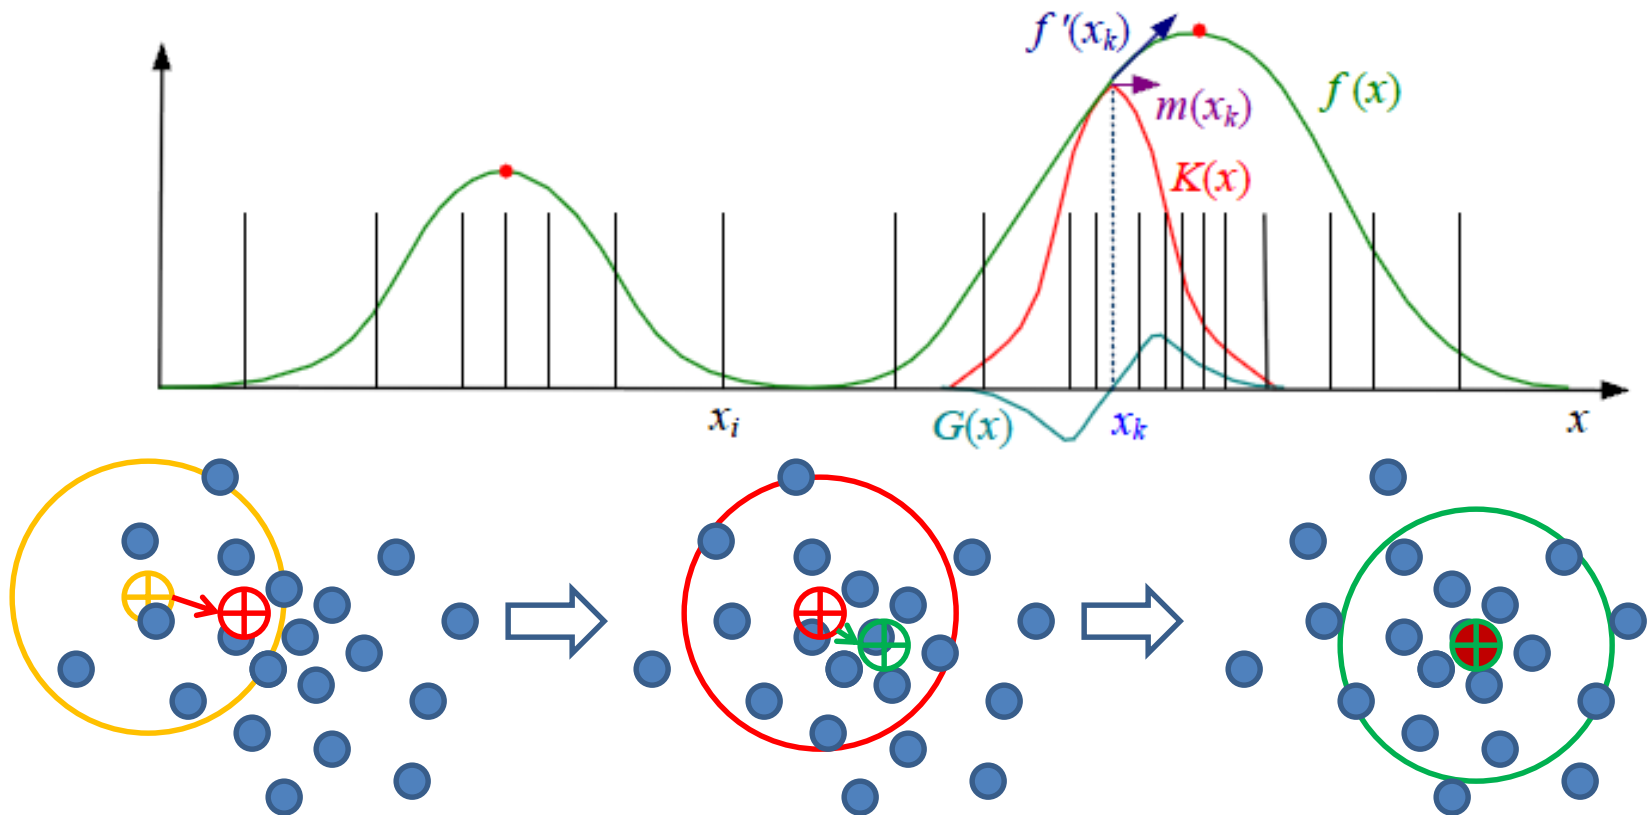
$$\begin{aligned} \nabla f(x) &= \sum_i (x_i - x) g\left(\frac{\|x - x_i\|^2}{h^2}\right) = \sum_i (x_i - x) G(x - x_i), g(r) = -k'(r) \Rightarrow \\ \nabla f(x) &= \left[\sum_i G(x - x_i) \right] m(x), G = \text{Derivative Kernel} \end{aligned}$$

where the vector $m(x) = \frac{\sum_i x_i G(x - x_i)}{\sum_i G(x - x_i)} - x$ is called **mean shift**, which is the difference between the weighted mean of the neighbours x_i around x and the current value of x

- The current estimate $y_{k+1} = y_k + m(y_k) = \frac{\sum_i x_i G(y_k - x_i)}{\sum_i G(y_k - x_i)}$

Mean Shift


- 1-D visualization of kernel density estimate $f(x)$, kernel K , its derivative G , and a mean shift m , weighted nearby points for re-estimation of the mean, red dots are local maxima to which mean shifts converge to




Mean Shift




(a)




(b)



(c)



(d)



(e)

Mean Shift

- Mean Shift for Visual Tracking
 - mean shift algorithm can be used for visual tracking
 - create a confidence map in the new image based on colour histogram of the object in the previous image
 - tracking target object in sequence by matching colour density
 - use mean shift to find the peak of a confidence map near the object's old position



References and Acknowledgements

- Chapter 3, 5 Szeliski 2010
- Some images drawn from Szeliski 2010 and web
- Other references:
 - T.F. Cootes and C.J. Taylor and D.H. Cooper and J. Graham. "Active shape models - their training and application". *CVIU*, 1995.
 - D. Comaniciu, V. Ramesh, and P. Meer. Real-time tracking of non-rigid objects using mean shift. CVPR, 2000